

HRIDYESH KUMAR

Machine Learning / AI Intern Candidate

Phone: +91-7302990253 | Email: hridyeshkumar03@gmail.com
LinkedIn: linkedin.com/in/hridyesh-kumar | GitHub: github.com/Hridye5h

SUMMARY

Final-year B.Tech student at IIIT Allahabad targeting Machine Learning / AI roles. Builds and rigorously evaluates models end-to-end — deep learning (PyTorch, CLIP, CNN ensembles) and gradient-boosted models — with a focus on honest evaluation (walk-forward validation, data-leakage audits). Also writes low-latency C++ (a limit-order-book matching engine at ~13M orders/s). Codeforces Expert, LeetCode Knight (top 6%).

TECHNICAL SKILLS

Languages: Python, C++, C, Java, SQL

ML & Deep Learning: PyTorch, Transformers, scikit-learn, LightGBM, NumPy, Pandas, OpenCV

ML Concepts: CNNs, Transformers, CLIP / Vision-Language Models, Multimodal Fusion, Gradient Boosting, Cross-Sectional Ranking, Walk-Forward Backtesting, Retrieval-Augmented Generation (RAG), Grad-CAM Explainability

Quant & Backend: Feature Engineering, Time-Series Validation, Risk Parity, FastAPI, Docker, Qdrant (Vector DB), Hugging Face, REST APIs

CS Fundamentals: Data Structures & Algorithms, Operating Systems, Object-Oriented Programming, Computer Networks, DBMS

PROJECTS

AlphaRank — ML Quant Trading System (NSE) | *Python, LightGBM, pandas, Flask* GitHub

- Built a systematic equity system that computes ~50 point-in-time features on ~380 NSE stocks and uses a multi-horizon (5/10/20-day) LightGBM model to rank stocks cross-sectionally each week, holding a risk-parity basket of the top 7 with volatility-adaptive stops, sector caps, and a drawdown kill-switch.
- Validated with a cost-adjusted walk-forward backtest (rolling retrain, horizon-length embargo); **caught and fixed a data-leakage bug** that had inflated the Information Coefficient from a fake 0.48 to an honest ~0.01–0.02, then rebuilt the metric as a per-date Spearman rank IC. Runs on free data in paper-trading mode with a live dashboard.

Facial Emotion Recognition — RAF-DB | *PyTorch, OpenCV, Grad-CAM, Gradio* GitHub | Live Demo

- Fine-tuned DDAMFN (a face-specialised attention CNN) on RAF-DB (7 emotions, ~15K images) and built a DDAMFN + ConvNeXt-V2 stacking ensemble reaching **91.6% test accuracy**, on par with published state-of-the-art.
- Added Grad-CAM explainability, deployed a live Gradio demo on Hugging Face Spaces, and fixed a data-leakage bug for an honest, leak-free train/validation/test split.

Multimodal Product Tagging — CLIP | *PyTorch, CLIP, Transformers, Gradio* GitHub | Live Demo

- Built a CLIP-based multimodal model that auto-tags products with 4 catalog attributes from image + title in one forward pass — **92% mean accuracy** (category 98%, gender 99%) across 30K products.
- Quantified robustness to missing metadata: with titles removed, mean accuracy holds at 82% as the model falls back on CLIP image features; deployed as a live Gradio demo.

Also: Match Engine — price-time-priority limit-order-book matching engine in C++ (~13M orders/s, p99 ~400ns; O(1) cancel, arena allocator, lock-free SPSC feed, 71 tests); **RepoSage** — agentic RAG over codebases with hybrid retrieval (dense + BM25, RRF), reaching MRR 0.81.

EDUCATION

Indian Institute of Information Technology, Allahabad	2023 – present
<i>B.Tech in Electronics and Communication Engineering</i>	<i>CGPA: 7.34</i>
St. Mary's Inter College, Etawah	2019 – 2022
<i>Class XII (CBSE) 94.4% — Class X (CBSE) 97.4%</i>	

CODING PROFILES

Codeforces: Expert (1701) — **LeetCode:** Knight (1863, top 6%) — **CodeChef:** 3-Star (1720)

ACHIEVEMENTS

Top 150 nationally — TenzorX 2026 National AI Hackathon (team).

NTSE Scholar (2020): Top 2000 out of 1.2 million applicants.

PRMO (2018): Top 456 in Uttar Pradesh — **VVM** (2020): Top 18 in Uttar Pradesh.